NL2code

Sai Krishna Karanam karanam.s@husky.neu.edu
Nischal Mahaveer Chand mahaveerchand.n@husky.neu.edu
Varun Sundar Rabindranath rabindranath.v@husky.neu.edu

I. Introduction

A. Importance of NL2code

Recent Integrated Development Environments (IDEs) provide a feature called "context-aware code completion" which helps programmers speed up the coding process. This is very helpful for programmers who are already aware of the programming language syntax and the design paradigms. But, for beginners, this can be time consuming.

The first step, for beginners, is to search how to implement the code in natural language. Due to the vast amount of resources available online, many unreliable or outdated, having a NL to code generator will make programming more enjoyable for beginners and more productive for veteran programmers. With the introduction of a new language almost everyday, it has become more and more important to establish an interface between NL and code.

We aim to create neural based models to generate code from NL. Below we describe several attempts of implemening various encoder-decoder models for the same. It should be noted that the project does not aim at providing a framework to build industry standard deployable code, but is rather a proof-of-concept. The scope of this project is limited to generate an accurate single line program statement given user intent in the form of a comment or description.

Examples

input	call the function sorted with argument x
ref.	sorted(x)
input	for every k in keys
ref.	for k in keys:

II. Related Word

Techniques to generate regular expressions [1], input parsers [2], UML diagrams, object oriented class layout, and general purpose code from Natural Language (NL) specification have been researched for a long time.

Recent approaches in converting NL to executable source code predominantly use neural network techniques. These approaches follow two basic patterns. One, is to directly convert NL to source code by modeling the task as a Sequence-to-Sequence conversion task, using Recurrent Neural Networks (RNN) (Lili Mou et al., 2015, Ling et al., 2016) [3], [4] to that end. The other is to generate Abstract Syntax Trees (ASTs) from NL input, then deterministically convert the produced AST to code (Pengcheng Yin et al., 2017) [5].

However, research [5] shows that the second approach is better than the first and emphasizes the use of syntax information of the target language, in NL to code tasks. We consider the recent research paper by Pengcheng Yin et. al. [5] , on syntactic neural model for code generation as the primary motivation for this project. [5] proposes an approach where they translate NL to an Abstract Syntax Tree (AST) first, using the grammar of the target language as a prior knowledge. They report a 10% absolute improvement in accuracy compared to the previous state-of-the-art on standard datasets. Since, converting from AST to code can be done deterministically, their system always produces syntactically correct, executable code. This aspect is lacking in the Sequence-to-Sequence models, where the correctness of the generated code is not guaranteed. They use a Bidirectional LSTM (BiLSTM) network for encoding each word of the NL input as a context specific embedding, and an RNN as a decoder to generate an output AST.

Maxim Robinovich et. al. [6], propose an approach similar to [5] for converting NL to AST. They confirm the experimental results of [5], thus again emphasizing the use of target language syntax. However Maxim Robinovich et. al., have evaluated their system on a wider range of datasets achieving near state-of-the-art results on average.

Pengcheg Yin et.al [5], report that 25% of the errors incurred by their system, on one of the datasets, is due to the generated code only partially implementing the required functionality. We hypothesize that this could be due to the standard datasets used for this task being noisy. We propose a denoising procedure, explained in the Data Preprocessing section to clean the dataset, thereby making the datasets more generic.

Barone et. al, 2017 [7], hereafter referred to as the "Edinburgh paper", explores the difficulty of the code generation task and critiques the currently available datasets for this task. They also comment that the effectiveness reported by [5] should not be taken at face value.

In the Edinburgh paper, Barone et al. scraped online open source Python repositories to extract "docstring-code" pairs. They propose that such data accurately represents the query/code pairs found in production environment. They used a Neural Machine Translation approach to translate docstring to code and reported a baseline BLEU score of 10.24, While [5] reports a BLEU score of 84.5 on the Django dataset.

We wanted to check how [5] performs on a similar dataset. We scraped query-code pairs from the same online open source Python repositories as [7]. [5], on our scraped

dataset reports a BLEU score of only 14.2 which reaffirms the [7]'s claim that the reported effectiveness of various models could not be taken at face value and that the task of code generation is more difficult than what is conceived by the community. Much of this misrepresentation could be attributed to the fact that the datasets currently available for this task are very limited and noisy.

III. Dataset

A. Standard Datasets

There are 3 datasets that have been used by the research community for the task of code-generation from the comments/query.

1) Django: Django dataset is a collection of 18000 lines of python code from the open source web framework, Django. The 18000 code statements have been manually annotated with a Natural language description. As a result of manual annotations, all the annotations are artificially descriptive of the corresponding code. This is the main drawback of the dataset, as the manual annotations rarely reflects the comments that programmers write in production environment. This can be seen in the following example.

```
comment:
  raise and exception InvalidCacheBackendError
  with string "Could not find config for '\%s'
  in settings.CACHES" as argument , replace '
  \%s ' with alias .

code:
  raise InvalidCacheBackendError ( "Could not
  find config for '\%s' in settings.CACHES" \%
  alias )
```

2) HearthStone: HearthStone dataset is a collection of 600 Python class - description pairs that implement cards for the online multiplayer card game HearthStone. The class - description pairs used to train and test are very similar, in that any model trained on this dataset would be able to generate only a single template of code, like the one presented in the example. We could not use this dataset as the scope of the dataset is very narrow and far from the case in general production environment.

```
card details:
  name: ['D', 'i', 'r', 'e', ' ',
        'W', 'o', 'l', 'f', ' ',
        'A', 'l', 'p', 'h', 'a']
  cost: ['2']
  type: ['Minion']
  rarity: ['Common']
  race: ['Beast']
  class: ['Neutral']
  description: ['Adjacent', 'minions', 'have',
        '+', '1', 'Attack', '.']
  health: ['2']
  attack: ['2']
  durability: ['-1']

correspoding code:
  class DireWolfAlpha(MinionCard):
    def __init__(self):
      super().__init__( "Dire Wolf Alpha", 2,
        CHARACTER_CLASS.ALL, CARD_RARITY.COMMON,
        minion_type=MINION_TYPE.BEAST)
    def create_minion(self, player):
      return Minion(2, 2, auras=[Aura(ChangeAttack(1),
                        MinionSelector(Adjacent()))])
```

3) Edinburgh: The Edinburgh dataset is a collection of 200k docstring-code pairs scraped from Opensource python repositories. The docstring-code pairs actually reflect the comments written in production environment. However, the docstring reflects what a particular entity (method, class, function etc.) represents rather than how it is implemented. For the task of generating source code, we need information on implemention logic thus making the dataset unusable for our task.

```
comment:
  Process elementNode by svgReader.

code:
  def processSVGElementtext(elementNode, svgReader):
    if svgReader.yAxisPointingUpward:
      return
    fontFamily = getStyleValue('Gentium Basic Regular',
        elementNode, 'font-family')
    fontSize = getRightStripAlphabetPercent(getStyleValue
        ('12.0', elementNode, 'font-size'))
    matrixSVG = getChainMatrixSVGIfNecessary(elementNode,
        svgReader.yAxisPointingUpward)
    loopLayer = svgReader.getLoopLayer()
    translate = euclidean.getComplexDefaultByDictionaryKeys(
        complex(), elementNode.attributes, 'x', 'y')
    for textComplexLoop in getTextComplexLoops(fontFamily,
      fontSize, elementNode.getTextContent(),
      svgReader.yAxisPointingUpward):
      translatedLoop = []
      for textComplexPoint in textComplexLoop:
        translatedLoop.append(textComplexPoint + translate)
      loopLayer.loops.append(matrixSVG.
        getTransformedPath(translatedLoop))
```

B. Our dataset

1) Scrapped Dataset: We scraped several of the Opensource python repositories from github and collected 12.5k pairs of code-comments after preprocessing the data. We found that the code is very project specific and the descriptions are context specific. This dataset also had the same issue as the Edinburgh dataset as in the descriptions mainly focused on why a code is written rather than how it's written.

2) Processed Django: Among all the datasets explored, the Django dataset represents the problem of converting NL to source code more accurately. But, as mentioned above the Django dataset is very noisy, hence we denoise it (techniques mentioned in data preprocessing) before using it.

Dataset Statistics	
Maximum comment length	133
Maximum code length	701
Average comment length	16.0
Average code length	8.7
Vocabulary size - Query	2205
Vocabulary size - Code	814
Vocabulary size - POS	44
Vocabulary size - Phrase	14

IV. Models

We describe four models, each is built upon Pengcheng's model, but encodes the input in a different method. First we start by briefly describing Pengcheng's model.

A. Pengcheng's model

Pengcheng recognized that adding syntax information to the model would give better result. Their model follows an encoder-decoder architecture, and takes the raw comment as input and generates an AST of the corresponding code as output.

The encoder comprises of an embedding layer and a BiLSTM layer. It takes a comment as input, embeds each word in the comment to give token embeddings TE_t, for each word t in the comment. Each TE_t is sequentially fed into the BiLSTM layer, to produces a Query Embedding (QE) of 128 dimensions. QE is passed to the decoder module.

The decoder uses an Recurrent Neural Network (RNN) to sequentially generate each node of the AST. Each node maps to a timestep in the RNN decoding process and thus, generating the AST can be interpreted as unrolling the RNN. The RNN also maintains an internal state to track the generation process at each timestep.

Pengcheng tackles this problem with a probabilistic grammar model of generating an AST y given NL description x: $p(y|x)$. The best possible AST $\hat{y}$ is given by:

$$\hat{y} = \arg\max_y p(y|x) \tag{1}$$

$\hat{y}$ is then deterministically converted to the corresponding Python code.

Figure 1: Architecture of Pengcheng's model. The encoder embeds each token of the NL input to give TE_t, which is then passed into the BiLSTM layer. The encoder produces a 128 dimension QE which is fed into the decoder.

B. Our models

As described, the decoder is already state-of-the-art, and needs no further modifications. All models described here-on use the same decoder architecture as Pengcheng's model with modifications only to the encoder and the input data. All models are trained and tested using the dataset described in Processed Django.

1) Basic Concat (BC): For our first attempt to incorporate syntax information into the encoder, we decided to concatenate (shown as ":") the POS ID and phrase ID of each token to the corresponding token embedding, giving the Augmented Token Embedding (ATE). The ATE is then fed into a modified BiLSTM layer that takes 130 dimension embeddings, rather than the default 128 dimensions.


Figure 2: Architecture of Basic Concat model. Notice that the POS ID and phrase ID are directly appended to the corresponding TE. The resulting 130 dimension embedding is passed to the BiLSTM.

TE_t dimension: 128
POS_ID_t and $Phrase\_ID_t$ dimension: 1 each; total 2
ATE_t dimension = 130

$ATE_t = [TE_t : POS\_ID_t : Phrase\_ID_t]$
$QE = BiLSTM(ATE)$

2) Linear Projection (LP): To add syntactic information over a sequence of tokens, we used an embedding layer for POS and phrase tags. The resulting TE, POS embedding (POSE), and Phrase embedding (PhE) are then concatenated to produce the ATE; which is a $(128 * 3)$ dimension vector. We then apply a linear projection (using a dense layer with the linear activation function) to give ATE Projected (ATEP) which is passed to BiLSTM.


Figure 3: Architecture of Linear Projection model. Each embedding - TE, POSE, and PhE are concatenated and projected (shown as purple circles). The resulting ATEP is passed to the BiLSTM.

TE_t dimension: 128
$POSE_t$ dimension: 128
PhE_t dimension: 128
Dense = $(128 * 3)$ input nodes, 128 output nodes

$ATE_t = [TE_t : POSE_t : PhE_t]$
$ATEP_t = \text{Dense}(ATE_t)$
$QE = \text{BiLSTM}(ATEP)$

3) Linear Projection Reduced Dimension (LPrd): Subsequently, we noticed that the POS and Phrase vocabulary sizes were relatively smaller than token vocabulary size. Thus, we reduce the embedding dimensions of POSE and PhE to 8 and 32 respectively. The process described in LP is then repeated.


Figure 4: Architecture of Linear Projection Reduced Dimension model. Same as LP, but with reduced dimensions for POSE and PhE.

TE_t dimension: 128
$POSE_t$ dimension: 8
PhE_t dimension: 32
Dense = (128 + 8 + 32) input nodes, 128 output nodes

$ATE_t = [TE_t : POSE_t : PhE_t]$
$ATEP_t = \text{Dense}(ATE_t)$
$QE = \text{BiLSTM}(ATEP)$

4) Raw Query Independent Preprojection (AdvLP): Rather than applying one linear projection on ATE, we apply two here, where the first is independent of the input query. POSE and PhE are concatenated and projected to create Augmentation Embedding (AE), which is then concatenated with TE and projected to produce QE. We hypothesize that this would preserve most of the information from the input query during projection.


Figure 5: Architecture of Raw Query Independent Preprojection model. POSE and PhE are concatenated and projected independent of the input query. The resulting is concatenated with TE and projected to produce QE.

TE_t dimension: 128
$POSE_t$ dimension: 128

PhE_t dimension: 128
Dense = (128 * 2) input nodes, 128 output nodes

$AE_t = \text{Dense}([POSE_t : PhE_t])$
$ATE_t = [TE_t : AE_t]$
$ATEP_t = \text{Dense}(ATE_t)$
$QE = \text{BiLSTM}(ATEP)$

V. Experiments

All decoder dimensions and configurations are left untouched and are thus the same as in Pengcheng's model. For each of the above models, the embedding sizes are as described in Models. Each model is run for a maximum of 50 epochs, with early stopping if the validation metrics do not change for 10 epochs.

VI. Results

We evaluate our model on two metrics: BLEU and Accuracy.
1) BLEU: NLTK's bleu_score function is used with NIST geometric sequence smoothing.
2) Accuracy: We use exact match for accuracy. +1 if the token matches the reference exactly, 0 otherwise. The total is normalized by the number of the prediction.

Models		Metrics	
		BLEU	Accu.
Pengcheng		84.5	71.6
Base		73.2	67.9
NL2code	BC	73.6	69.4
	LP	74.3	69.7
	LPrd	73.6	69.0
	AdvLP	73.7	69.1

VII. Conclusion

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like "Huardest gefburn"? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

References

[1] J. Berant, A. Chou, R. Frostig, and P. Liang, "Semantic parsing on freebase from question-answer pairs," in Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing, 2013, pp. 1533–1544.

[2] T. Lei, F. Long, R. Barzilay, and M. Rinard, "From natural language specifications to program input parsers," in Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), vol. 1, 2013, pp. 1294–1303.

[3] L. Mou, R. Men, G. Li, L. Zhang, and Z. Jin, "On end-to-end program generation from user intention by deep neural networks," arXiv preprint arXiv:1510.07211, 2015.

[4] W. Ling, E. Grefenstette, K. M. Hermann, T. Kočiskỳ, A. Senior, F. Wang, and P. Blunsom, "Latent predictor networks for code generation," arXiv preprint arXiv:1603.06744, 2016.

[5] P. Yin and G. Neubig, "A syntactic neural model for general-purpose code generation," arXiv preprint arXiv:1704.01696, 2017.

[6] M. Rabinovich, M. Stern, and D. Klein, "Abstract syntax networks for code generation and semantic parsing," arXiv preprint arXiv:1704.07535, 2017.

[7] A. V. M. Barone and R. Sennrich, "A parallel corpus of python functions and documentation strings for automated code documentation and code generation," arXiv preprint arXiv:1707.02275, 2017.